

# Machine Learning

## Lecture 3.2: $k$ -Nearest Neighbors

Marc Sebban and Amaury Habrard

LABORATOIRE HUBERT CURIEN, UMR CNRS 5516  
Université Jean Monnet Saint-Étienne


# Bayesian Error

## Bayesian Error

The bayesian error  $\epsilon^*$  is the lowest possible error rate (or irreducible error) for any hypothesis  $h$ .

$$\epsilon^* = \int_{x \in R_i \text{ s.t. } y \neq y_i} p(y_i|x)p(x)dx$$

where  $x$  is an instance,  $y$  its corresponding label,  $R_i$  is the area/region that a classifier function  $h$  classifies as  $y_i$ .



# Bayesian Classifier

## Bayesian Classifier

The Bayesian classifier predicts the optimal class  $y^* \in \mathcal{Y}$  given an example  $\mathbf{x} \in \mathcal{X}$  by applying the Maximum a posteriori (MAP) decision rule:

$$\forall y_j \in \mathcal{Y}, p(y_j|\mathbf{x}) = \frac{p(\mathbf{x}|y_j) \cdot p(y_j)}{p(\mathbf{x})}$$

$$y^*(\mathbf{x}) = \arg \max_c p(y_c|\mathbf{x}).$$

that corresponds to  $y^*(\mathbf{x}) = \arg \max_c p(\mathbf{x}|y_c) \cdot p(y_c)$

If this calculation is possible, the Bayesian classifier is optimal from a probabilistic point of view, with an associated error  $\epsilon^*$ .

# Underlying conditions to solve this problem

To compute  $\epsilon^*$ , one needs some priors:

- 1 Know the *a priori* probabilities  $p(y_j)$  of the different classes.
- 2 Know the probabilities of the observations given the classes  $p(\mathbf{x}|y_j)$ .

Unfortunately,  $p(y_j)$  and  $p(\mathbf{x}|y_j)$  are unknown. One needs to estimate these two quantities from the training sample  $S$ .

# Estimation of the a priori probability of the classes $p(y_j)$

What about  $p(y_j)$ ?

An unbiased estimate of  $p(y_j)$  is given by the observed frequency

$\hat{p}(y_j) = \frac{|S_j|}{|S|}$  where  $|S_j|$  is the number of training examples belonging to the class  $y_j$ .

# Estimation of the conditional probabilities $p(\mathbf{x}|y_j)$

What about  $p(\mathbf{x}|y_j)$ ?

We can distinguish two types of approaches:

- 1 The **parametric methods** which assume that  $p(\mathbf{x}|y_j)$  follows a given statistical distribution. In this case, the problem to solve consists in estimating the parameters of the considered distribution (e.g. normal distribution with  $\mu$  and  $\sigma$  or Binomial distribution with  $p$ ).
- 2 The **non parametric methods** which do not impose any constraint about the underlying distribution, and for which the densities  $p(\mathbf{x}|y_j)$  are locally estimated around  $\mathbf{x}$ .

# Non parametric methods

- No assumption is made on the underlying distribution...
- ... we only suppose that this target distribution is **locally regular**.
- The objective is to estimate  $p(\mathbf{x}|y_j)$ . Since this must be done  $\forall y_j \in \mathcal{Y}$ , for the sake of simplicity, let us focus on  $p(\mathbf{x})$  first.

# Non parametric methods

- Let  $p$  be the unknown target probability density. The probability  $\mathcal{P}$  to observe  $\mathbf{x}$  in a volume  $V$  is:

$$\mathcal{P} = \int_V p(\mathbf{x}) d\mathbf{x}$$

- Assuming that  $p(\mathbf{x})$  does not significantly change in  $V$  (**locally regular**), we can approximate  $\mathcal{P}$  such that:

$$\hat{\mathcal{P}} \simeq p(\mathbf{x}) \times V \quad (1)$$

- $\mathcal{P}$  can also be estimated by the proportion of training data in  $V$ :

$$\hat{\mathcal{P}} \simeq \frac{k}{n} \quad (2)$$

- Therefore, we can deduce that

$$p(\mathbf{x}) \approx \frac{k}{nV} = \hat{p}(\mathbf{x}).$$



## Theorem

Let us denote by  $\hat{p}_n(\mathbf{x}) = \frac{k_n}{nV_n}$  the estimate of  $p(\mathbf{x})$  from a training sample  $S$  of size  $n$ . When  $n$  is increasing,  $\hat{p}_n(\mathbf{x})$  converges to  $p(\mathbf{x})$  if the following three conditions are fulfilled:

$$\lim_{n \rightarrow \infty} V_n = 0$$

$$\lim_{n \rightarrow \infty} k_n = \infty$$

$$\lim_{n \rightarrow \infty} \frac{k_n}{n} = 0$$

## Interpretation


If  $n$  is high, we get:

- a good estimate  $\hat{p}(\mathbf{x})$  (resp.  $\hat{p}(\mathbf{x}|y_j)$  if we take into account the class) of  $p(\mathbf{x})$  (resp.  $p(\mathbf{x}|y_j)$ ).
- a good approximation of the Bayesian method.

# Non parametric methods

## $k$ -Nearest Neighbors (Cover & Hart 1968)

It turns out that the  $k$ -nearest neighbor method satisfies the previous conditions: **fixing a number  $k_n$  of examples** (growing w.r.t.  $n$ ) and **adapting a volume  $V_n$**  (e.g. a hypersphere centered at  $\mathbf{x}$ ) such that  $k_n$  examples are in the volume.




# $k$ -nearest neighbors

## $k$ NN as a good way to estimate $p(\mathbf{x})$



To estimate  $p(\mathbf{x})$  from the  $n$  training examples of  $S$ , let us build a hypersphere centered at  $\mathbf{x}$  that contains  $k_n$  examples.

$\hat{p}_n(\mathbf{x}) = \frac{k_n/n}{V_n}$  is a good estimate of  $p(\mathbf{x})$





# Example 1 with $k = \sqrt{n}$

$$p(\mathbf{x}) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}\mathbf{x}^2}$$



# Example 1 with $k = \sqrt{n}$


$$p(\mathbf{x}) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}\mathbf{x}^2}$$



## $k$ -NN as a classifier

Let us suppose we have  $n_j$  examples in  $S$  of class  $y_j$ , such that  $\sum_j n_j = n$ . Suppose the hypersphere contains  $k_j$  examples of class  $y_j$  ( $\sum_j k_j = k$ ). If we aim at classifying a new example  $\mathbf{x}$  with its  $k$  nearest-neighbors, and applying the Bayes rule, we get:

$$h(\mathbf{x}) = \arg \max_j \frac{\hat{p}(\mathbf{x}|y_j) \cdot \hat{p}(y_j)}{\hat{p}(\mathbf{x})} = \arg \max_j \frac{\frac{k_j}{V_n \times n_j} \times \frac{n_j}{n}}{\frac{k}{n \times V_n}} = \arg \max_j \frac{k_j}{k}$$



# $k$ -nearest neighbors algorithm

How to classify an unknown example  $\mathbf{x}$  using  $S$ ?

**Input:**  $\mathbf{x}, S, d$

**Output:** class of  $\mathbf{x}$

**foreach**  $(\mathbf{x}', y') \in S$  **do**

    | Compute the distance  $d(\mathbf{x}', \mathbf{x})$ ;

Sort the  $n$  distances by increasing order;

Count the number of occurrences of each class  $y_j$  among the  $k$  nearest neighbors;

Assign to  $\mathbf{x}$  the most frequent class;

# Special case of the $(k = 1)$ -nearest neighbor rule

## Convergence properties of the $(k = 1)$ -nearest neighbor rule

### Theorem

Let  $\mathbf{x}'$  be the nearest neighbor of  $\mathbf{x}$ ,

$$\lim_{n \rightarrow \infty} p(d(\mathbf{x}, \mathbf{x}') > \epsilon) = 0, \forall \epsilon > 0$$

### Corrolary

If  $n \rightarrow \infty$ ,  $p(y_j|\mathbf{x}') \approx p(y_j|\mathbf{x})$ .



# Special case of the ( $k = 1$ )-nearest neighbor rule

## Proof.

Let  $\mathcal{P}$  be the probability that the hypersphere  $s(\mathbf{x}, \epsilon)$  centered at  $\mathbf{x}$  of radius  $\epsilon$  does not contain any point of  $S$ :

$$\begin{aligned}\mathcal{P} &= p(\mathbf{x}_1 \notin s(\mathbf{x}, \epsilon), \dots, \mathbf{x}_n \notin s(\mathbf{x}, \epsilon)) = \prod_{i=1}^n p(\mathbf{x}_i \notin s(\mathbf{x}, \epsilon)) \\ &= \prod_{i=1}^n (1 - p(\mathbf{x}_i \in s(\mathbf{x}, \epsilon))) = (1 - p_\epsilon)^n.\end{aligned}$$

Since the nearest neighbor  $\mathbf{x}'$  of  $\mathbf{x}$  is also outside of the sphere, we get

$$p(d(\mathbf{x}, \mathbf{x}') > \epsilon) = (1 - p_\epsilon)^n$$

$$\text{therefore, } \lim_{n \rightarrow \infty} p(d(\mathbf{x}, \mathbf{x}') > \epsilon) = \lim_{n \rightarrow \infty} (1 - p_\epsilon)^n = 0$$



## Theorem

*The generalization error  $\epsilon_{1NN}$  of the 1-nearest neighbor rule is bounded by twice the (optimal) bayesian error  $\epsilon^*$ .*

$$\epsilon_{1NN} \leq 2\epsilon^*$$

## Corrolary

Half of the information about the true class of an example  $\mathbf{x}$  is contained in its nearest neighbor  $\mathbf{x}'$ .

# Special case of the ( $k = 1$ )-nearest neighbor rule

Proof.

Bayesian error:

$$\begin{aligned}\forall \mathbf{x} \in \mathcal{X}, \epsilon^*(\mathbf{x}) &= \min \{p(y_1|\mathbf{x}), p(y_2|\mathbf{x})\} \\ &= p(y_{\min}|\mathbf{x})\end{aligned}$$

1NN error:

$$\begin{aligned}\epsilon_{1NN}(\mathbf{x}) &= p(y_1|\mathbf{x})p(y_2|\mathbf{x}') + p(y_2|\mathbf{x})p(y_1|\mathbf{x}') \\ (\text{if } n \text{ is large}) &\approx p(y_1|\mathbf{x})p(y_2|\mathbf{x}) + p(y_2|\mathbf{x})p(y_1|\mathbf{x}) \\ &= 2p(y_{\min}|\mathbf{x})(1 - p(y_{\min}|\mathbf{x})) \leq 2p(y_{\min}|\mathbf{x}) = 2\epsilon^*(\mathbf{x})\end{aligned}$$



We can deduce that:

In the discrete case

$$\epsilon^* = \sum_{\mathbf{x} \in \mathcal{X}} p(y_{min}|\mathbf{x}) \cdot p(\mathbf{x})$$

$$\epsilon_{1NN} = \sum_{\mathbf{x} \in \mathcal{X}} (p(y_1|\mathbf{x})p(y_2|\mathbf{x}') + p(y_2|\mathbf{x})p(y_1|\mathbf{x}')) \cdot p(\mathbf{x}) \leq 2\epsilon^*$$


In the continuous case

$$\epsilon^* = \int_{\mathbf{x} \in \mathcal{X}} f(y_{min}|\mathbf{x}) \cdot f(\mathbf{x}) d\mathbf{x}$$

$$\epsilon_{1NN} = \int_{\mathbf{x} \in \mathcal{X}} (f(y_1|\mathbf{x})f(y_2|\mathbf{x}') + f(y_2|\mathbf{x})f(y_1|\mathbf{x}')) \cdot f(\mathbf{x}) d\mathbf{x} \leq 2\epsilon^*$$


# Graphically

The property  $\epsilon^* \leq \epsilon_{1NN} \leq 2\epsilon^*$  can be graphically illustrated



# Effect of $k$ on the estimation quality

$$\epsilon^* \leq \epsilon_{kNN} \leq \epsilon_{(k-1)NN} \leq \dots \leq \epsilon_{1NN} \leq 2\epsilon^*$$



The previous property holds only asymptotically  $\rightarrow$  possible compromise between  $k$  and  $n$ :  $k = \sqrt{\frac{n}{|Y|}}$ .

# Problems of $k$ -NN

- To converge, the  $k$ -nearest neighbor algorithm requires a large number of training examples.
- However, a large number of training examples implies a large space/time complexity.

Two strategies to overcome these problems:

- Reduce the size of  $S$  while keeping the most relevant examples (e.g. the condensed nearest neighbor rule (Hart 1968)).
- Simplify the calculation of the nearest-neighbor.

# Data reduction techniques

*Preliminary step:* remove from  $S$  the outliers and the examples of the bayesian error region.

**Input:**  $S$

**Output:**  $S_{cleaned}$

Split randomly  $S$  into two subsets  $S_1$  and  $S_2$ ;

**while** *no stabilization of  $S_1$  and  $S_2$*  **do**

    Classify  $S_1$  with  $S_2$  using the 1-NN rule;

    Remove from  $S_1$  the misclassified instances;

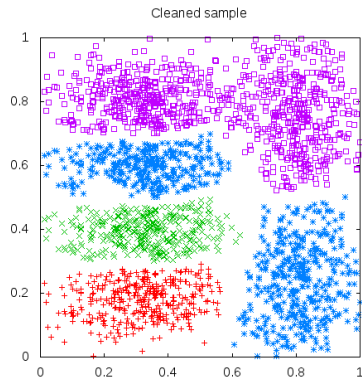

    Classify  $S_2$  with the new set  $S_1$  using the 1-NN rule;

    Remove from  $S_2$  the misclassified instances;

$S_{cleaned} = S_1 \cup S_2$ ;



# Illustration



# The condensed nearest neighbor rule (CNN)

*Second step:* remove the irrelevant examples.

**Input:**  $S$

**Output:** STORAGE

$\text{STORAGE} \leftarrow \emptyset$  ;  $\text{DUSTBIN} \leftarrow \emptyset$ ;

Draw randomly a training example from  $S$  and put it in STORAGE;

**while** *no stabilization of STORAGE* **do**

**foreach**  $\mathbf{x}_i \in S$  **do**

**if**  $\mathbf{x}_i$  *is correctly classified with STORAGE using the 1-NN rule*

**then**

$\text{DUSTBIN} \leftarrow \mathbf{x}_i$


**else**

$\text{STORAGE} \leftarrow \mathbf{x}_i$


**return** STORAGE;

# Illustration

Cleaned sample



Condensed sample




# How to speed-up the nearest-neighbor calculation?

Most of the approaches are based on the **triangle inequality property**


$$\forall(\mathbf{x}, \mathbf{y}, \mathbf{z}) \in \mathcal{X}^3, d(\mathbf{x}, \mathbf{z}) \leq d(\mathbf{x}, \mathbf{y}) + d(\mathbf{y}, \mathbf{z}).$$

- Let  $\mathbf{x}$  be the example to classify by the NN rule. Let  $\mathbf{y}$  be the current NN of  $\mathbf{x}$  at a distance  $\delta$ .
- Let  $\mathbf{z}$  be the next example. If  $d(\mathbf{x}, \mathbf{z}) \leq \delta$  then update the current NN. Otherwise, remove the following examples:
  - in the sphere centered at  $\mathbf{z}$  and of radius  $d(\mathbf{x}, \mathbf{z}) - \delta$ ,
  - out of the sphere centered at  $\mathbf{z}$  and of radius  $d(\mathbf{x}, \mathbf{z}) + \delta$ ,



# How to speed-up the nearest-neighbor calculation?

- $d(\mathbf{v}_1, \mathbf{z}) \leq d(\mathbf{x}, \mathbf{z}) - \delta \Rightarrow d(\mathbf{v}_1, \mathbf{z}) + \delta \leq d(\mathbf{x}, \mathbf{z})$  (1).
- By the triangle inequality, we know that  $d(\mathbf{x}, \mathbf{z}) \leq d(\mathbf{x}, \mathbf{v}_1) + d(\mathbf{v}_1, \mathbf{z})$  (2).
- From (1) and (2) we get  $d(\mathbf{v}_1, \mathbf{z}) + \delta \leq d(\mathbf{x}, \mathbf{v}_1) + d(\mathbf{v}_1, \mathbf{z}) \Rightarrow \delta \leq d(\mathbf{x}, \mathbf{v}_1)$
- Therefore,  $\mathbf{v}_1$  cannot be the NN of  $\mathbf{x}$ . The same proof can be used for  $\mathbf{v}_2$ .



## Some methods about fast kNN algorithms

- In 2D or 3D: graph-based searching methods such as Voronoi diagram and proximity graph.
- In higher spaces: ball-trees, kd-trees, metric-trees, quadtree, R-trees, k-Means-k-NN.

# Conclusion

- With a sufficiently large number of training examples, a NN classifier is able to converge towards very complex target functions.
- It is simple and theoretically well founded.
- There exist several solutions to overcome its algorithmic complexity issues (time and space).